

let's summarize what the **Dissertation Analysis** solution does today, then map each of its major components to the appropriate layer in the Spanda Platform (Platform Layer, Domain Layer, Solutions/App Layer). Finally, we'll show how the **0.5 “platformed”** version of the Dissertation Analysis integrates into a single-node deployment, paving the way for the future multi-node fabric in v1.0.

1. High-Level Purpose of Dissertation Analysis

The **Dissertation Analysis** repository is an **end-to-end prototype** (or standalone solution) for:

1. **Ingesting Academic Documents** (theses, dissertations, reports)
2. **Text Preprocessing & Chunking** (splitting long documents into sections/chapters)
3. **AI-Driven Analysis**
 - **Semantic Understanding:** Summaries, key points, or concept extraction
 - **Rubric Alignment / Feedback:** Mapping textual segments to rubric criteria and providing automated or semi-automated feedback
 - **Plagiarism / Similarity Checks:** (If configured) Identifying suspiciously similar text or references
4. **Reporting & Grading Suggestions**
 - Generating **targeted feedback** for instructors or TAs
 - Possibly outputting a grade “draft” or recommended rubric scores
5. **(Optional) Web or CLI Interface**
 - Depending on the user's environment, some combination of Python scripts, notebooks, or minimal front-end might be present to upload a dissertation file and see the analysis results

In short, it addresses the **academic workflow** of evaluating a dissertation's structure, originality, and alignment with standardized requirements or rubrics.

2. Current Repo Structure & Components

Although the exact folder layout may differ, the **Dissertation Analysis** repo typically includes (or is likely to include) components such as:

1. **Data Ingestion & Processing**
 - Scripts to read PDF/Word documents (e.g., `ingest.py`, `parser.py`, etc.)
 - Utilities for chunking large text into smaller sections, extracting metadata, or cleaning up references.
2. **LLM Integration / Model Scripts**
 - Code that calls out to large language models (e.g., via Hugging Face Transformers, GPT-based endpoints, or local LLM runners like Ollama).

- Fine-tuned model weights or references to external model hosts.
- 3. **Rubric & Feedback Logic**
 - A set of **rules** or python modules that compare a chunk of text against specific rubric criteria (organization, clarity, novelty, referencing, etc.).
 - Possibly a “scoring” or “ranking” script that compiles final feedback.
- 4. **User Interface / CLI**
 - A minimal web UI (Flask, FastAPI, or Streamlit) or a Jupyter notebook that demonstrates how a user might upload a dissertation, run an analysis, and view results.
 - Alternatively, a command-line approach with a `main.py` script that processes an input file and prints out suggestions or a grading breakdown.
- 5. **Config & Dependencies**
 - Requirements files, environment configs, or partial Dockerfiles that define how everything runs.
- 6. **(Optional) Database Adapters**
 - Some repos might store results or logs in a local SQLite/MySQL.
 - Others might have a simple “no-database” approach with everything in memory.

Note: Over time, these pieces will be migrated into dedicated microservices or modules so they integrate nicely into the Spanda.AI multi-layer ecosystem.

3. Mapping to the Spanda.AI Platform Layers

3.1 Platform Layer

- **Model Serving Infrastructure:** If the Dissertation Analysis is using local LLM inference (e.g., Ollama or any container-based serving mechanism), that “inference runtime” belongs to the **Platform Layer**.
- **Data & Message Brokers:** Any references to Kafka, MySQL, Redis, or Prometheus monitoring in the repo will ultimately be “lifted” and placed in the **Platform**. For instance:
 - A container that provides **MySQL** for storing user data or feedback results.
 - **Redis** for quick caching of chunked text or partial inference results.
 - **Kafka** for processing large volumes of text in an asynchronous pipeline.
 - **Prometheus** or “DockProm” for capturing usage metrics and system performance.

Rationale: The “foundation” for compute, data management, and logging/monitoring is a **Platform** responsibility.

3.2 Domain Layer (EdTech)

- **Dissertation-Specific Logic:** The Python scripts or classes that define how to chunk, parse, and interpret dissertations.

- **Rubric / Grading Modules:** Code that compares text sections to domain rubrics.
- **Model Fine-Tuning:** If you’ve built or integrated custom language model weights specifically tuned for academic or dissertation text, that belongs in the **EdTech Domain** subfolder.
- **Agent Orchestration (LangGraph):** In 0.5, if we add an agent workflow that’s specifically curated for dissertation tasks (chain-of-thought logic or multi-step parsing for academic text), it would primarily be in the **Domain Layer**, although the agent framework’s libraries (LangGraph) get installed in the Platform. The “domain flow” or prompt templates, however, live in EdTech code.

Rationale: Domain Layer = all the specialized business or sector-specific knowledge. In EdTech’s case, it’s the dissertation analysis methodology and the academic rubrics themselves.

3.3 Solutions Layer (“Dissertation Analysis App”)

- **Presentation & APIs:** Any front-end code, UI dashboards, or REST endpoints that present final grading suggestions to a user (instructor, TA, admin).
- **Integration Adapters:** If the solution needs to plug into an LMS (Learning Management System), or to push results to a school’s administrative platform, that bridging code sits in the Solutions Layer.
- **Reporting & Export:** PDF or HTML generation for a final “feedback” report that an instructor might share with a student.

Rationale: This top layer is all about end-user experiences, integration points, and “last-mile” delivery of the domain’s AI capabilities.

4. The 0.5 “Platformed” Version of Dissertation Analysis

With **Spanda.AI v0.5**, we’re taking the functionality from the standalone Dissertation Analysis repo and **containerizing** it into a **single-node** (Mac or PC, CPU or GPU) setup. Here’s how that works in practice:

1. **docker-compose.yml (Platform)**
 - The base `docker-compose.yml` references containers like Kafka, MySQL, Redis, Prometheus, etc.
 - Also includes a container for Ollama or other local model servers.
 - Optionally includes a container for the agent framework (LangGraph) if it’s integrated in 0.5.
2. **EdTech Domain Container**
 - A container image that packages up the domain logic for dissertation analysis: chunking scripts, rubric logic, any specialized LLM-fine-tuning code.
 - This container depends on the Platform containers (for data storage, logs, GPU runtime if available).
3. **Dissertation Analysis App Container (Solutions Layer)**

- A minimal web or API-based front-end that an instructor or user sees.
 - Connects to the domain container (via internal Docker network) to run the analysis pipeline on uploaded files.
 - 4. **Single-Machine Advantage**
 - Everything runs on your local workstation or server using `docker-compose up`.
 - Great for demonstrations, PoCs, and small-scale usage.
 - Supports GPU acceleration if you have a compatible GPU. If not, everything just runs on CPU.
 - 5. **Agent Flows**
 - If **LangGraph** is included:
 - The dissertation text can be orchestrated through multiple LLM calls (e.g., summarization → rubric check → feedback generation).
 - This pipeline is *still domain-specific* but leverages the shared Platform capabilities (logging, concurrency, HPC resources, etc.).
-

5. Future Transition to the v1.0 “Fabric” Deployment

While **v0.5** is single-node, the **v1.0** fabric approach will enable:

- **Distributed Nodes:** Multiple CPU/GPU nodes—local, cloud, or hybrid—joined into a single orchestrated environment.
 - **Regional Deployment:** US-based nodes for one part of the institution, APAC-based nodes for another, all coordinating seamlessly.
 - **Scalable Dissertation Analysis:** The same domain container can auto-scale horizontally if thousands of dissertations need parallel analysis.
 - **Cross-Domain Expansion:** HRTech or SportsTech containers can join the same fabric, each specialized in different tasks but sharing the underlying Platform.
-

6. Final Takeaways

1. **Today’s Repo:**
 - A self-contained application that ingests academic papers, runs AI-based analysis, and outputs structured feedback.
 - Bundles domain logic, partial data/ML code, and UI or demonstration interfaces into one place.
2. **Where Each Piece Lands:**
 - **Platform Layer:** Kafka, Zookeeper, MySQL, Redis, Prometheus, Ollama (model serving), plus agent libraries (LangGraph) at a system level.
 - **Domain Layer:** EdTech’s specialized chunking, rubrics, LLM fine-tuning, and academic heuristics.

- **Solutions Layer:** The final “Dissertation App” front-end, APIs, and integration points.
- 3. **0.5 “Platformed”:**
 - Single-node, Docker Compose–based deployment that merges these pieces into containers.
 - Demonstrates GPU or CPU readiness, local box usage, and how an agent workflow can orchestrate multiple LLM calls for dissertations.
- 4. **Looking Ahead to 1.0:**
 - A **fabric-based** architecture that scales horizontally or globally, orchestrating multi-region nodes for advanced performance, resilience, and multi-domain synergy.

By **migrating** the key elements of the **Dissertation-Analysis** repo into their rightful places in the Spanda.AI **Platform**, **Domain**, and **Solutions** layers, we ensure that the **EdTech** offering (dissertation analysis) becomes more modular, scalable, and easier to integrate with future domain expansions—ultimately delivering a robust, enterprise-ready AI ecosystem.